

Brittany Hancock  
IT 420: Business Intelligence  
Instructor: Lenore Montalbano  
OLTP vs OLAP Client Consulting Report

**Introduction:** OLTP systems handle daily transactions and store detailed data that changes often. These databases capture raw information used later by OLAP systems. OLAP, on the other hand, gathers that transactional data in batches and organizes it for analysis. While OLTP focuses on accuracy and real-time updates, OLAP is designed for quick access when creating reports, dashboards, and charts.

**Design differences:** A normalized OLTP design keeps each piece of data in its own field, for instance, separating a customer's first and last name. OLAP systems use a denormalized approach, combining related details to make reports easier to build. Instead of calculating totals each time a query runs, OLAP stores results like total sales directly in the table, which saves time later. In contrast, OLTP must calculate those figures during each query, slowing performance as the database grows.

OLAP structures are built for large-scale analysis, while OLTP systems ensure reliable transactions. In a star schema, OLAP saves processed data for reuse, and each record has a surrogate key that allows quick, indexed joins between tables. For example, a CustomerKey in the dimension table connects directly to a matching CustomerKey in the fact table. A fully normalized OLTP system does not use this type of relationship.

**Performance comparison:** See Figures 1 through 14 in the Appendix for query results, ETL validation, and schema performance examples.

**Business risks:** Running analytical queries on an OLTP database can slow it down because it must repeatedly group and calculate results. These systems rely on multiple joins and aggregations, whereas OLAP stores summarized and indexed data ready for use. As the amount of information grows, OLTP performance drops and can disrupt regular operations.

**Recommendation:** OLAP offers a more efficient structure for analytics, reporting, and visualization. It supports high-volume summaries and scales smoothly as data expands. Surrogate keys make relationships between tables faster to query. Because OLAP organizes and pre-calculates data in advance, it provides accurate, dependable results for business planning and decision making.

## Appendix

Figure 1. SSIS Control Flow Overview: displays the full ETL process from staging to warehouse, including data flow execution and truncation steps for both environments.

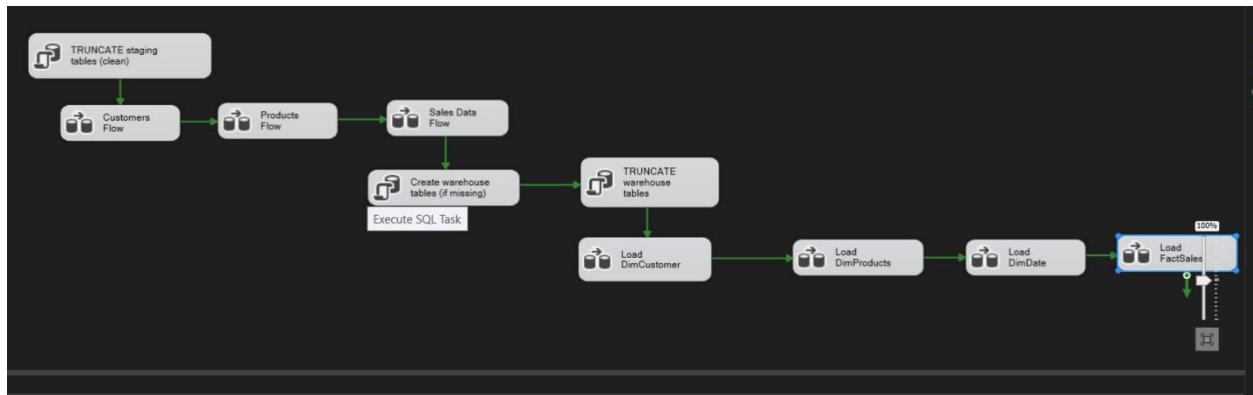


Figure 2. SSIS Data Flow For FactSales: shows the transformation pipeline where surrogate keys are looked up, derived columns are calculated, and clean data is loaded into the FactSales table.

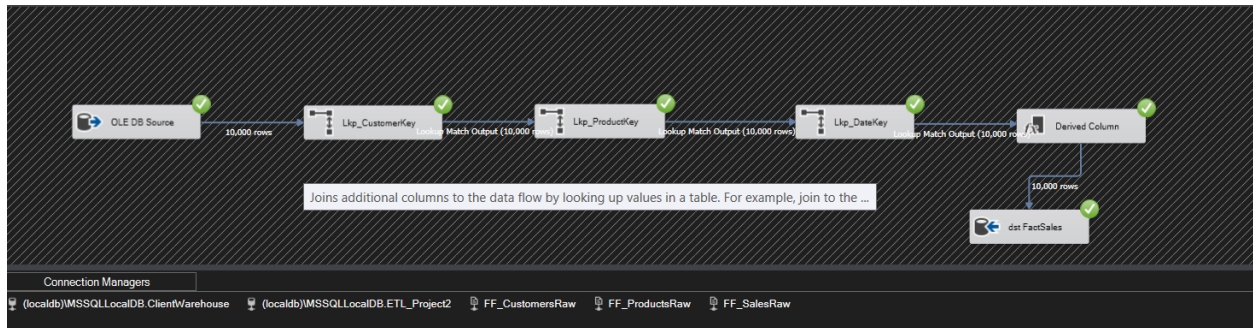
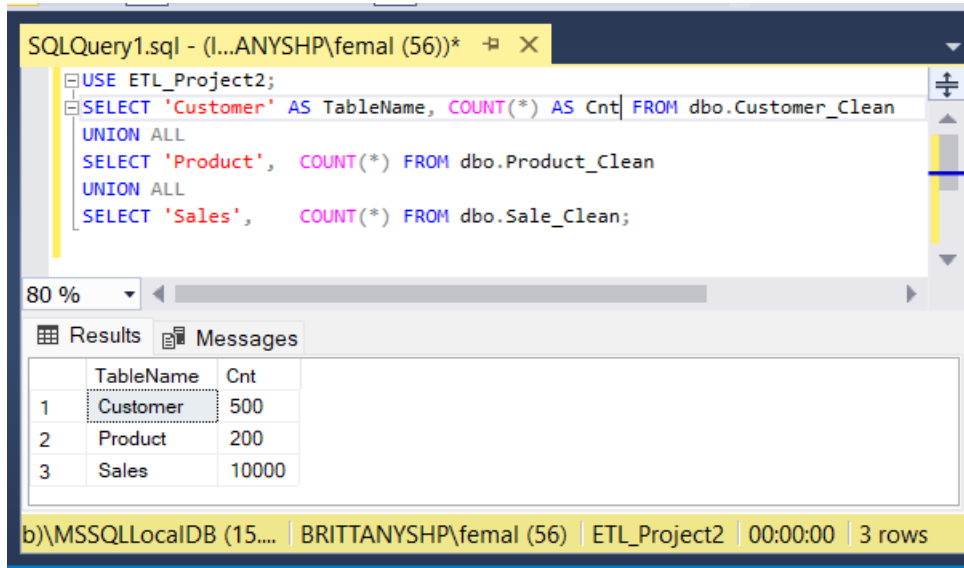


Figure 3. SSIS Package Successful Execution: indicates the ETL workflow ran successfully from start to finish with all data flow tasks completed.



Figure 4. Staging Database Row Counts (ETL\_Project2): row counts from the cleaned staging tables confirm successful data import: 500 customers, 200 products, and 10,000 sales records.



The screenshot shows a SQL query window titled 'SQLQuery1.sql - (\\...\\ANYSH\\femal (56))'. The query is as follows:

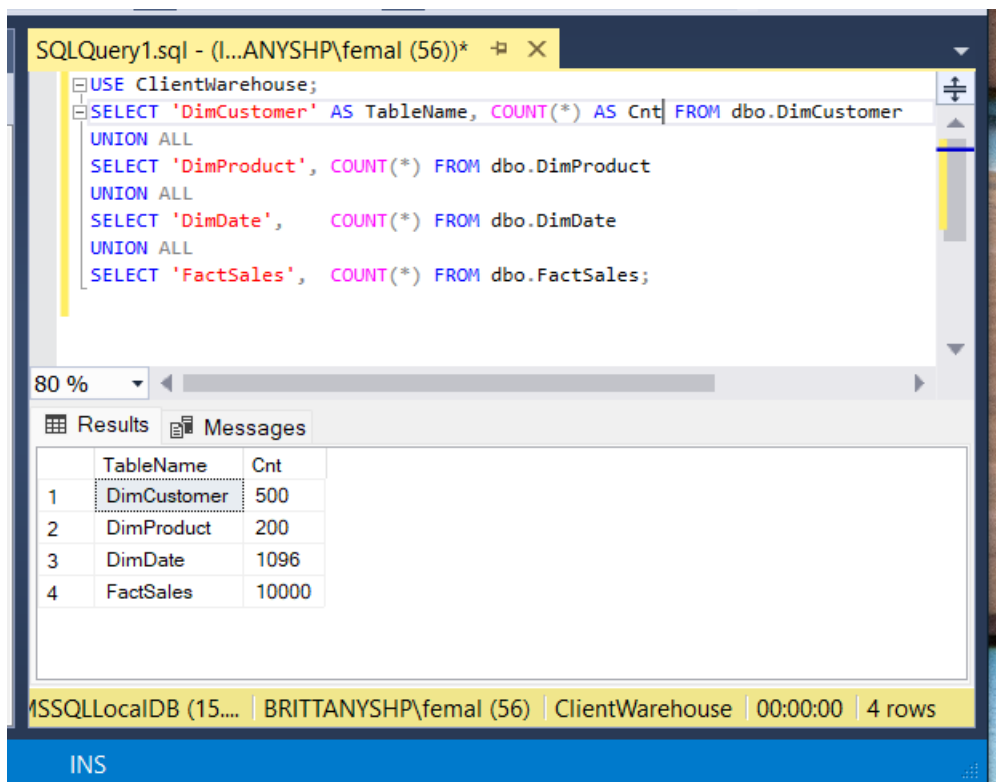
```
USE ETL_Project2;
SELECT 'Customer' AS TableName, COUNT(*) AS Cnt FROM dbo.Customer_Clean
UNION ALL
SELECT 'Product', COUNT(*) FROM dbo.Product_Clean
UNION ALL
SELECT 'Sales', COUNT(*) FROM dbo.Sale_Clean;
```

The results grid displays the following data:

	TableName	Cnt
1	Customer	500
2	Product	200
3	Sales	10000

The status bar at the bottom indicates: 'b)\\MSSQLLocalDB (15... | BRITTANYSH\\femal (56) | ETL\_Project2 | 00:00:00 | 3 rows'.

Figure 5. Warehouse Database Row Counts (ClientWarehouse): row counts in dimension and fact tables confirm successful ETL transfer and matching totals across both environments.



The screenshot shows a SQL query window titled 'SQLQuery1.sql - (\\...\\ANYSH\\femal (56))'. The query is as follows:

```
USE ClientWarehouse;
SELECT 'DimCustomer' AS TableName, COUNT(*) AS Cnt FROM dbo.DimCustomer
UNION ALL
SELECT 'DimProduct', COUNT(*) FROM dbo.DimProduct
UNION ALL
SELECT 'DimDate', COUNT(*) FROM dbo.DimDate
UNION ALL
SELECT 'FactSales', COUNT(*) FROM dbo.FactSales;
```

The results grid displays the following data:

	TableName	Cnt
1	DimCustomer	500
2	DimProduct	200
3	DimDate	1096
4	FactSales	10000

The status bar at the bottom indicates: 'MSSQLLocalDB (15... | BRITTANYSH\\femal (56) | ClientWarehouse | 00:00:00 | 4 rows'.

INS

Figure 6. TotalSaleAmount Validation Query: verification query checks consistency between computed totals and stored TotalSaleAmount values, confirming data accuracy (no mismatches found).

The screenshot displays the SQL Server Enterprise Manager interface. The top pane shows a query window titled "SQLQuery1.sql - (I...ANYSHP\femal (56))\*" containing the following SQL code:

```
USE ClientWarehouse;  
SELECT TOP 25  
    SalesKey, Quantity, UnitPrice, TotalSaleAmount,  
    CAST(Quantity * UnitPrice AS DECIMAL(12,2)) AS ComputedTotal  
FROM dbo.FactSales  
WHERE ABS(TotalSaleAmount - (Quantity * UnitPrice)) > 0.01;  
-- (No rows = good)
```

Below the query window, the "Results" tab is active, showing a table with the following columns: SalesKey, Quantity, UnitPrice, TotalSaleAmount, and ComputedTotal. The table is currently empty.

The bottom status bar indicates the connection is to "MSSQLLocalDB (15.... | BRITTANYSHIP\femal (56) | ClientWarehouse" and shows "00:00:00 | 0 rows".

Figure 7. Referential Integrity Validation Query: tests for missing Customer, Product, and Date keys in FactSales, confirming referential integrity with zero missing values.

The screenshot displays the SQL Server Enterprise Manager interface. The top pane shows a query window titled 'SQLQuery1.sql - (\\...\\ANYSH\\femal (56))\*'. The query contains three sections, each testing for missing keys in the FactSales table:

```
USE ClientWarehouse;

-- Missing Customers
SELECT COUNT(*) AS MissingCustomers
FROM dbo.FactSales f
LEFT JOIN dbo.DimCustomer c ON f.CustomerKey = c.CustomerKey
WHERE c.CustomerKey IS NULL;

-- Missing Products
SELECT COUNT(*) AS MissingProducts
FROM dbo.FactSales f
LEFT JOIN dbo.DimProduct p ON f.ProductKey = p.ProductKey
WHERE p.ProductKey IS NULL;

-- Missing Dates
SELECT COUNT(*) AS MissingDates
FROM dbo.FactSales f
LEFT JOIN dbo.DimDate d ON f.DateKey = d.DateKey
WHERE d.DateKey IS NULL;
```

The bottom pane shows the 'Results' tab with three tables, each containing one row with the value 0:

MissingCustomers	
1	0

MissingProducts	
1	0

MissingDates	
1	0

The status bar at the bottom indicates the connection is 'MSSQLLocalDB (15.... | BRITTANYSH\\femal (56) | ClientWarehouse | 00:00:00 | 3 rows'. The bottom-most bar shows 'INS'.

Figure 8. OLTP Query - Total Revenue by Customer: example of OLTP query aggregating total revenue by customer in the transactional database (ETL\_Project2).

SQLQuery1.sql - (I...ANYSHP\femal (56))\*

```

USE ETL_Project2;
SELECT
    c.CustomerName,
    SUM(s.Quantity * s.Price) AS TotalRevenue
FROM dbo.Sale_Clean s
JOIN dbo.Customer_Clean c ON s.CustomerID = c.CustomerID
GROUP BY c.CustomerName
ORDER BY TotalRevenue DESC;

```

80 %

Results Messages Execution plan

	CustomerName	TotalRevenue
1	Michelle Brown	97483.00
2	Terry Wheeler	73434.00
3	Alexis Schmitt	67728.00
4	Ashley Davis	67668.00
5	Amber Harris	67140.00
6	Maria Pierce	66794.00
7	Jose Williams	65232.00
8	David Reyes	63084.00
9	Robert Nelson	62854.00
10	Carrie Brady	62294.00
11	Samuel Stevens	61820.00
12	Monica Vasquez	61057.00
13	Emily Gonzalez	61045.00
14	John Armstrong	60826.00

\MSSQLLocalDB (15... | BRITTANYSHP\femal (56) | ETL\_Project2 | 00:00:00 | 499 rows

Figure 9. OLAP Query - Total Revenue by Customer: equivalent OLAP query aggregating pre-calculated TotalSaleAmount by customer in the warehouse (ClientWarehouse).

SQLQuery1.sql - (I...ANYSHP\femal (56))\*

```

USE ClientWarehouse;
SELECT
    dc.CustomerName,
    SUM(fs.TotalSaleAmount) AS TotalRevenue
FROM dbo.FactSales fs
JOIN dbo.DimCustomer dc ON fs.CustomerKey = dc.CustomerKey
GROUP BY dc.CustomerName
ORDER BY TotalRevenue DESC;

```

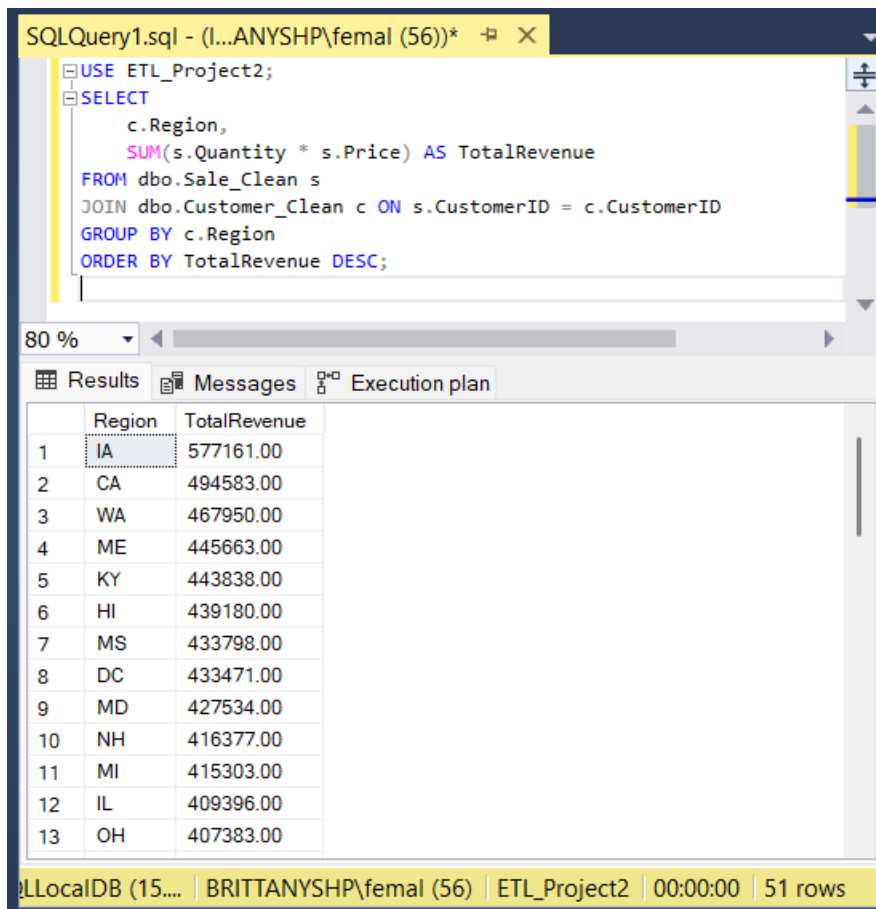
80 %

Results Messages Execution plan

	CustomerName	TotalRevenue
1	Michelle Brown	97483.00
2	Terry Wheeler	73434.00
3	Alexis Schmitt	67728.00
4	Ashley Davis	67668.00
5	Amber Harris	67140.00
6	Maria Pierce	66794.00
7	Jose Williams	65232.00
8	David Reyes	63084.00
9	Robert Nelson	62854.00
10	Carrie Brady	62294.00
11	Samuel Stevens	61820.00
12	Monica Vasquez	61057.00
13	Emily Gonzalez	61045.00

IDB (15.... | BRITTANYSHP\femal (56) | ClientWarehouse | 00:00:00 | 499 rows

Figure 10. OLTP Query - Total Revenue by Region: demonstrates regional revenue analysis on normalized OLTP data, requiring live aggregation and joins.



The screenshot displays the SQL Server Enterprise Manager interface. The top pane shows a query window titled 'SQLQuery1.sql - (\\...\\ANYSHF\\femal (56))\*'. The query is as follows:

```
USE ETL_Project2;
SELECT
    c.Region,
    SUM(s.Quantity * s.Price) AS TotalRevenue
FROM dbo.Sale_Clean s
JOIN dbo.Customer_Clean c ON s.CustomerID = c.CustomerID
GROUP BY c.Region
ORDER BY TotalRevenue DESC;
```

The bottom pane shows the 'Results' tab with a table containing 13 rows of data. The table has two columns: 'Region' and 'TotalRevenue'. The data is sorted in descending order of TotalRevenue.

	Region	TotalRevenue
1	IA	577161.00
2	CA	494583.00
3	WA	467950.00
4	ME	445663.00
5	KY	443838.00
6	HI	439180.00
7	MS	433798.00
8	DC	433471.00
9	MD	427534.00
10	NH	416377.00
11	MI	415303.00
12	IL	409396.00
13	OH	407383.00

The status bar at the bottom indicates the connection is to 'LocalDB (15.... | BRITTANYSHF\\femal (56) | ETL\_Project2' and shows '00:00:00 | 51 rows'.



Figure 11. OLAP Query - Total Revenue by Region: aggregates pre-calculated TotalSaleAmount by region in the warehouse (ClientWarehouse), showing faster performance compared to OLTP query.

The screenshot displays the SQL Server Enterprise Manager interface. The top pane shows a query window titled 'SQLQuery1.sql - (L...ANYSHP\femal (56))\*'. The query is as follows:

```
USE ClientWarehouse;
SELECT
    dc.Region,
    SUM(fs.TotalSaleAmount) AS TotalRevenue
FROM dbo.FactSales fs
JOIN dbo.DimCustomer dc ON fs.CustomerKey = dc.CustomerKey
GROUP BY dc.Region
ORDER BY TotalRevenue DESC;
```

Below the query window, the 'Results' tab is active, showing a table with 13 rows. The columns are 'Region' and 'TotalRevenue'. The data is sorted by TotalRevenue in descending order. The status bar at the bottom indicates 'calDB (15.... | BRITTANYSHP\femal (56) | ClientWarehouse | 00:00:00 | 51 rows'.

	Region	TotalRevenue
1	IA	577161.00
2	CA	494583.00
3	WA	467950.00
4	ME	445663.00
5	KY	443838.00
6	HI	439180.00
7	MS	433798.00
8	DC	433471.00
9	MD	427534.00
10	NH	416377.00
11	MI	415303.00
12	IL	409396.00
13	OH	407383.00

Figure 12. OLTP Query - Total Revenue by Year: displays yearly revenue totals computed in real time from the transactional database (ETL\_Project2) using live aggregation.

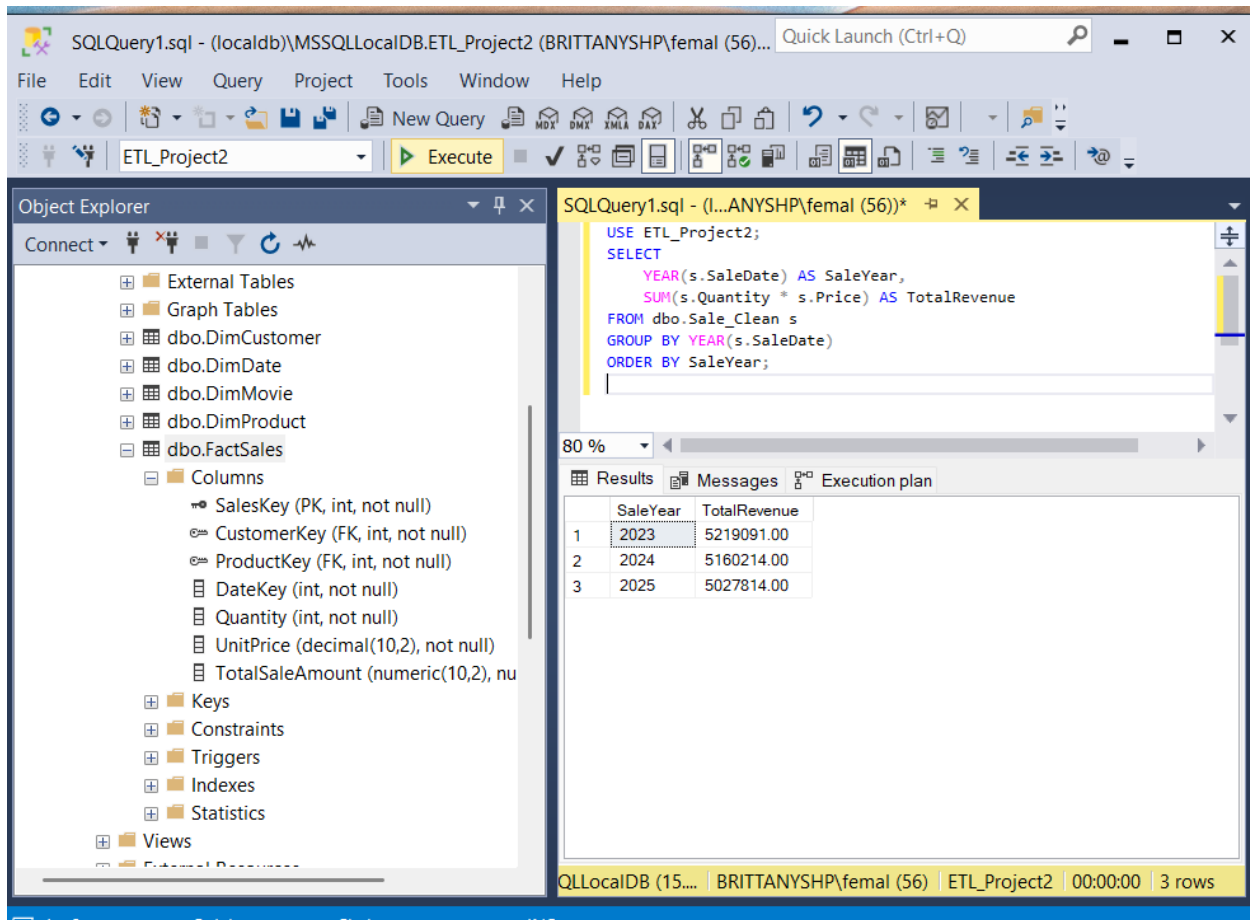


Figure 13. OLAP Query - Total Revenue by Year: shows the same yearly revenue analysis executed in the warehouse (ClientWarehouse), using pre-aggregated data for optimized query speed.

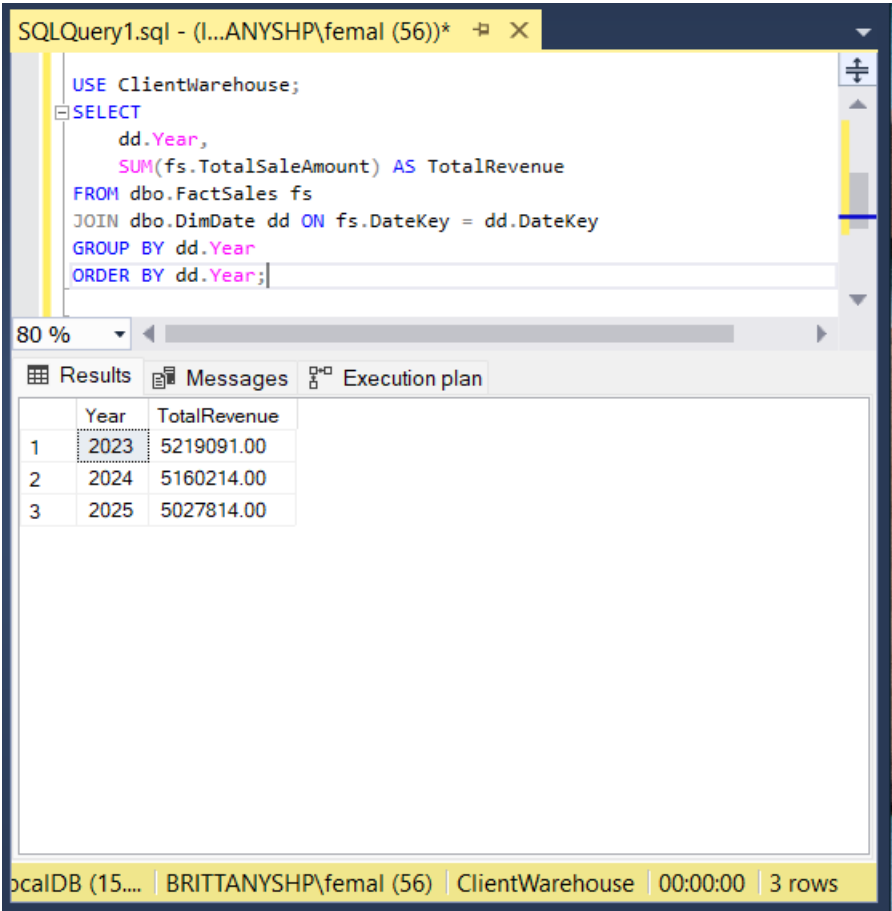


Figure 14. SSMS Schema Comparison: side-by-side view of the OLTP (ETL\_Project2) and OLAP (ClientWarehouse) database structures, showing normalized staging tables versus denormalized warehouse dimensions and fact table.

